论坛

General Discussion

General Discussion

What's your reading list? Share with us!

Yiwei Chen Mentor General Discussion · 2 years ago

Hello!

I believe people have some deep learning blogs/articles/papers that you find them useful or interesting! Do you like to share with us?

Some of mentors compiled a quick list just as a bit reference. If you like to recommend some other DL materials, you are always welcome and don't hesitate to reply! :D

=== Directly related to the course ===

Dropout

Srivastava, N., Hinton, G., Krizhevsky, A., Sutskever, I., and Salakhutdinov, R. (2014). Dropout: A simple way to prevent neural networks from overfitting. Journal of Machine Learning Research, 15 , 1929–1958.

Initialization

- He, K., Zhang, X., Ren, S., and Sun, J. (2015). Delving deep into rectifiers: Surpassing human-level performance on ImageNet classification. arXiv preprint arXiv:1502.01852 .
- Glorot, X. and Bengio, Y. (2010). Understanding the difficulty of training deep feedforward neural networks. In AISTATS'2010 .

Optimization

- RMSprop: http://www.cs.toronto.edu/~tijmen/csc321/slides/lecture_slides_lec6.pdf
- Kingma, D. and Ba, J. (2014). Adam: A method for stochastic optimization. arXiv preprint arXiv:1412.6980 .
- Dauphin, Y., Pascanu, R., Gulcehre, C., Cho, K., Ganguli, S., and Bengio, Y. (2014). Identifying and attacking the saddle point problem in high-dimensional non-convex optimization. In NIPS'2014 .

Hyperparameter tuning

- Bergstra, J. and Bengio, Y. (2012). Random search for hyper-parameter optimization. J. Machine Learning Res., 13 , 281–305.
- Bergstra, J, et. al. Algorithms for Hyper-Parameter Optimization. Advances in Neural Information Processing Systems (pp. 2546-2554).

Batch norm

- Ioffe, S. and Szegedy, C. (2015). Batch normalization: Accelerating deep network training by reducing internal covariate shift.

=== Additional materials ===

- Kaggle tutorials using examples for MNIST
- This online book - which makes use of javascript to make things intuitive
- Youtube channels - Siraj Raval , Two Minute Papers, etc
- Applying dropout units for image recognition competition: Going Deeper with Convolutions
- fast.ai - Making neural nets uncool again

33 个赞 Reply 关注此讨论

最早 最热门 最新

KW Karl Wiklund · a year ago

I'm already a working professional, and am doing this course as much to improve my python as to learn about some of the more "modern" techniques that have become popular since I graduated.

My list is:

Simon Haykin "Neural Networks and Learning Machines"

Simon Haykin "Adaptive Filter Theory" (Also Sayed's book, which is more current these days)

James Candy "Bayesian Signal Processing"

MacKay "Information Theory, Inference and Learning Algorithms"

Kevin Murphy "Machine Learning: A Probabilistic Perspective"

Stephen Boyd "Convex Optimization"

Courville, et al "Deep Learning"

Also, these text windows are super-small! It's really hard to type!

3 个赞 回复

Anurag Katakkar · a year ago

Hi!

One of the things that I find most useful to keep up with recent developments in the tech world is, surprisingly, Twitter.

My Twitter is entirely dedicated to following only tech, and more specifically, Machine Learning related Twitter accounts.

It is very, very interesting to see great researchers like Ian Goodfellow share their thoughts on contemporary research and also the best research/deep learning practices.

Though I'm quite new to Twitter myself, I've found these threads extremely interesting. Hope you enjoy this mode of learning as much as I do!

1 赞 回复

J. Christopher Bare · 2 years ago

Martin Gorner has a nice set of materials called *TensorFlow and Deep Learning without a PhD.* Slides, video, and code-lab:

https://docs.google.com/presentation/d/1TVixw6ItiZ8igjp6U17tcgoFrLSaHWQmMOwjlgQY9co/pub

https://www.youtube.com/watch?v=vq2nnJ4g6N0

https://codelabs.developers.google.com/codelabs/cloud-tensorflow-mnist/#0

0 个赞 回复

KC分钟

Kelvin Chan Mentor · 2 years ago

I think in addition to papers, it will be great if we also put a note on such list if the authors had open-src it and written in frameworks (tensorflow, Keras, etc).

0 个赞 回复

Paulo Abelha Ferreira Mentor · 2 years ago

Hi everyone! :)

Everything from David MacKay. I just love the way he explains and connect things.

To go deeper, the usual ML bibles. This Harvard course website has also some great pointers:

https://www.seas.harvard.edu/courses/cs281/

Best,

2 个赞 回复

Zohair Shafi Mentor · 2 years ago

Here are some more interesting papers and articles:

Ian Goodfellow's Generative Adversarial Models : https://arxiv.org/abs/1406.2661

Geoffrey Hinton's Capsule Networks (Article Explanation) : Part 1, Part 2, Part 3

I highly recommend visiting https://distill.pub. It presents research papers in a very easy to understand and graphic manner and explains them, as opposed to traditional PDF's.

Hope you like these! Happy learning!

2 个赞 回复

Joe Fernandez · 2 years ago

"Deep Learning" by Goodfellow, Bengio, and Courville, is a real good soup-to-nuts book on NNs. There's a free online version @ http://www.deeplearningbook.org/ and it also has links to some of Goodfellow's lecture videos.

4 个赞 回复

Jess Topple · 2 years ago

Here's a book recommendation: Hands-On Machine Learning with Scikit-Learn and TensorFlow

3 个赞 回复

Jess Topple · 2 years ago · 已编辑

coursera

Here's something that I hope will be interesting to people in this course: I came up with an optimization method insp[illegible]ol. I think it might be novel and would love to know what people here think of it!

I wrote a little paper about it and posted my code for anyone to try here: https://www.coursera.org/learn/deep-neural-network/discussions/weeks/2/threads/YOna-LiBEee[illegible]

您计算机的时区似乎与 Coursera 帐户的时区设置(America/Los_Angeles) 不符。

更改 Coursera 时区设置

0 个赞 隐藏 4 回复

Ricardo Cruz · 2 years ago

Hi Jess, that looks interesting, thanks for sharing. Anyone interested, please continue this discussion here. Let's keep this thread about literature.

0 个赞

Jaganath Babu · 2 years ago

Does any one have the downloaded notebook (without your codes) for 1 course assignments ? I have completed my first course (Deep learning and NN) last month. But, when i had some doubts and tried to access my assignments, All assignments were locked. If someone could hep me out, it could be of much help. Note: I am not asking for your codes, I just want the notebook with default codes that were present before starting the assignment. It would be of much help, if anyone can send me the notebook templates to jgun220@gmail.com. Thanks !!!

0 个赞

分钟

Yiwei Chen Mentor · 2 years ago

Hi Jaganath,

Please refer to this post, or the pinned post "FAQs and Tips for the Assignment Notebooks". That post provides a solution when being locked as well as other information such as downloading. I believe you can find your answer there :)

Since your reply is not related to the topic of this thread, I'm afraid I need to delete it later

Regards,

Yiwei

0 个赞

DW Darryl Wright · 8 months ago

I've created a reading list based on the interview video with Yoshua Bengio. You can access it here:

https://dr-darryl-wright.github.io/reading/list/2019/03/21/yoshua-bengio-interview-reading-list.html

I haven't gotten through them all yet, I tried to select the papers that best cover the topics mentioned in the video. Hope its helpful!

Thanks,

Darryl.

1 赞

Y 回复

回复

1

Y 回复

回复